

Generative Adversarial Networks: A Review


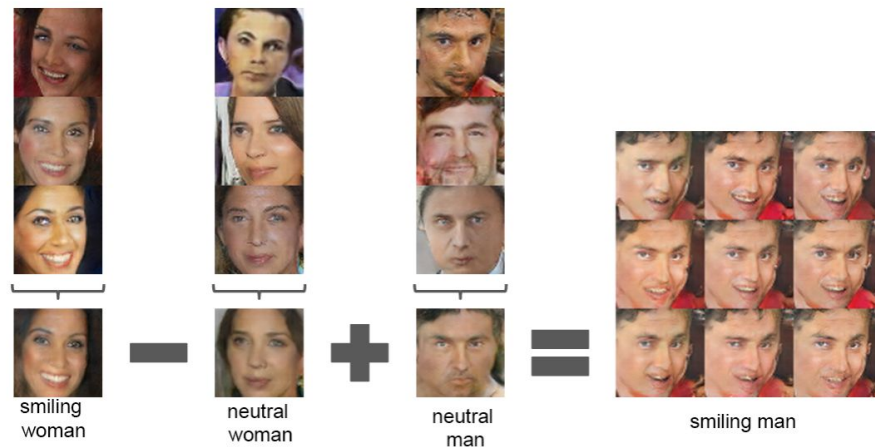
CSIC 5011: Topological and Geometrical Data Reduction and Visualization

Spring 2020

Jose Vinicius de Miranda Cardoso, Yixin Men, Shunkang Zhang

Applications

- High-resolution Image Synthesis, Image Inpainting, Image Editing, Image to image translation
- Object Detection
- Anomaly Detection in Medical Field
- Music and Video Generation
- Attention Prediction
- Financial Time-series Modeling



Introduction

- Generative adversarial networks (GANs) and other adversarial methods are based on a game-theoretical perspective on joint optimization of two neural networks as players in a game.
- Generative Adversarial Networks (GAN) have received wide attention in the machine learning field for their potential to learn high-dimensional, complex real data distribution.
- In contrast to discriminative models, *generative models* aim to learn the underlying distribution of the data and the generative process that creates them.
- D is trained to maximize the probability of assigning the correct label to both the training samples and samples generated by G. Simultaneously, G is trained to minimize the $\log(1-D(G(z)))$.

GAN Structure

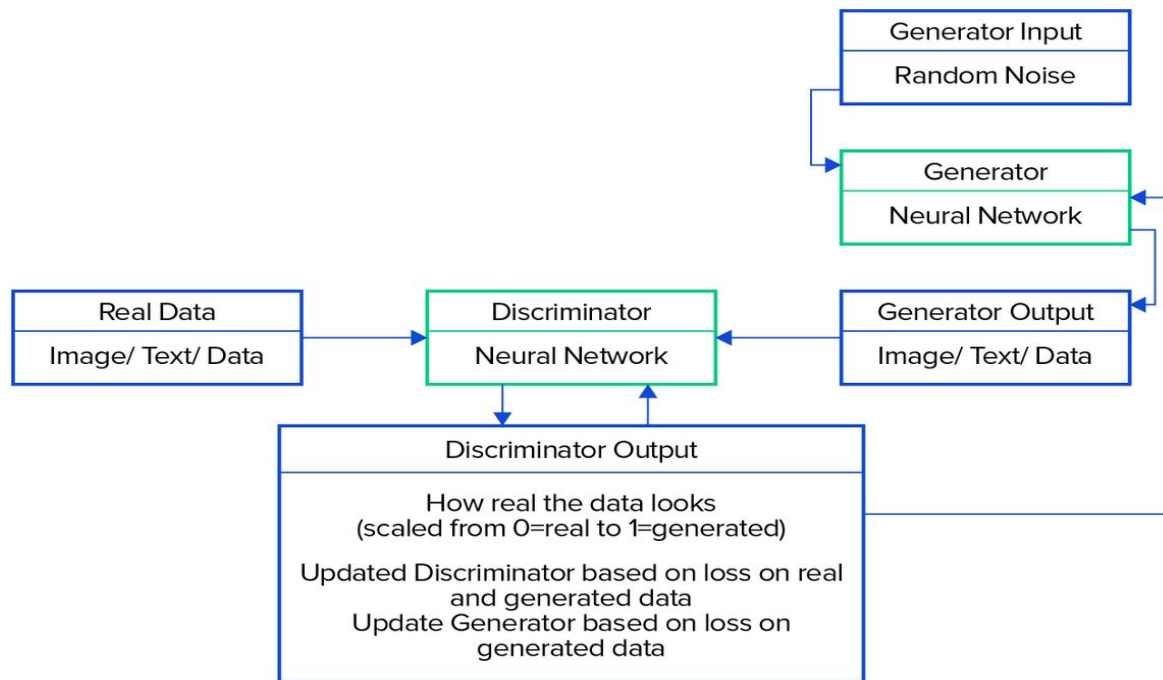


Figure 1. Generative adversarial network. The generator G takes a noise vector z sampled from a distribution p_z as input and uses fully connected or convolutional layers to transform this vector into a sample x . The discriminator D tries to distinguish these samples from samples drawn from the real data distribution p_{data} .

Generative adversarial networks consist of two neural networks.

The first network, the *generator*, tries to generate synthetic but perceptually convincing samples $\mathbf{x} \in p_{\text{fake}}$ that appear to have been drawn from a real data distribution p_{data} . It transforms noise vectors \mathbf{z} drawn from a distribution p_z into new samples, i.e., $\mathbf{x} = G(\mathbf{z})$.

The second network, the *discriminator*, has access to real samples from p_{data} and to the samples generated by G , and tries to discriminate between these two. GANs are trained by solving the following optimization problem that the discriminator is trying to maximize and the generator is trying to minimize:

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z} [\log(1 - D(G(\mathbf{z})))] ,$$

where G is the generator, D is the discriminator, $V(D, G)$ is the objective function, p_{data} is the distribution of real samples, and p_z is a distribution from which noise vectors are drawn, e.g., a uniform distribution or spherical Gaussian. The final layer of the discriminator network contains a sigmoid activation function, so that $D(\mathbf{x}), D(G(\mathbf{z})) \in [0, 1]$.

Basic training algorithm

Algorithm 1: Minibatch stochastic gradient descent training of generative adversarial nets.

1 **for** *number of training iterations* **do**

2 ▷ Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.

3 ▷ Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.

4 ▷ Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

5 ▷ Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.

6 ▷ Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

7 **end**

Experiments with synthetic data

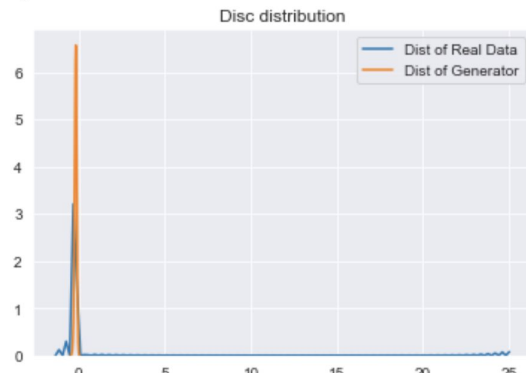
We generate data under Huber Contamination model:

$$x_i \sim 0.8\mathcal{N}(0, I_p) + 0.2\mathcal{N}(5 * 1_p, 2 * I_p),$$


with $N = 50000$ samples and dimension $p = 50$.

Experiments with synthetic data


Epoch:25, LossD/G:1.2810/-0.6130



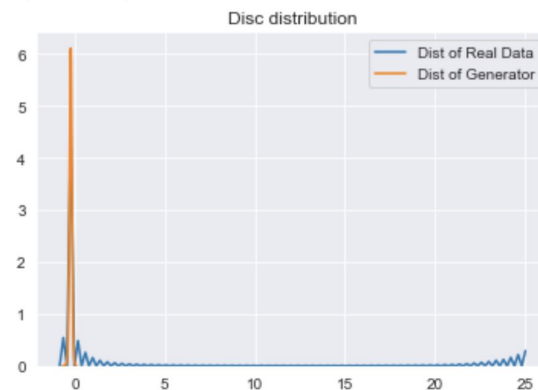
Epoch:50, LossD/G:1.2371/-0.5869



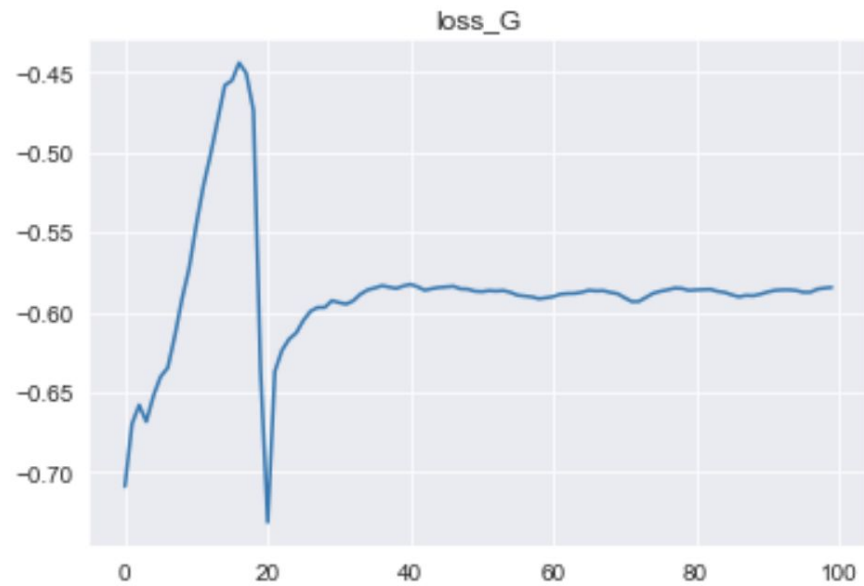
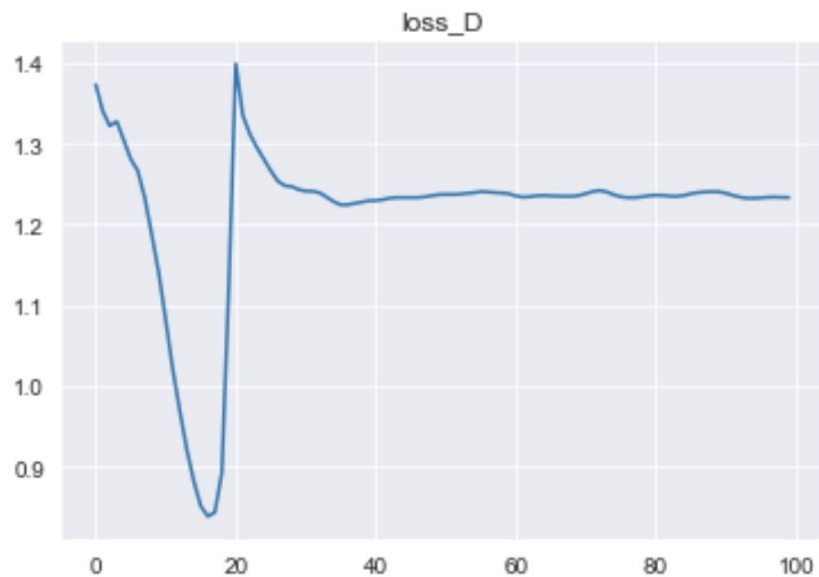
Epoch:75, LossD/G:1.2368/-0.5882




Epoch:100, LossD/G:1.2335/-0.5845




Experiments with synthetic data



Experiments with synthetic data: outlier detection

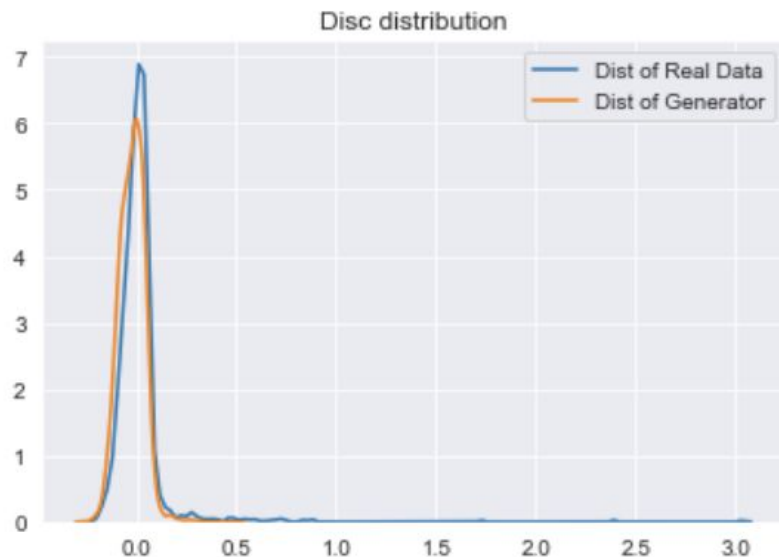


Stock data experiment




Stock data experiment

Epoch:250, LossD/G:1.3699/-0.6800



Epoch:1000, LossD/G:1.3104/-0.6430



More applications in stock markets

- **CorrGAN: Sampling Realistic Financial Correlation Matrices Using Generative Adversarial Networks**
<https://arxiv.org/abs/1910.09504>
- **Quant GANs: Deep Generation of Financial Time Series**
<https://arxiv.org/abs/1907.06673>

Reformulate GAN in function space

Given $\mu \in \mathcal{P}(\mathbb{R}^d)$ with the density q , we use the f -divergence to measure the discrepancy between μ and ν which is defined as

$$\mathbb{D}_f(q\|p) = \int p(\mathbf{x}) f\left(\frac{q(\mathbf{x})}{p(\mathbf{x})}\right) d\mathbf{x}, \quad (1)$$

where $f : \mathbb{R}^+ \rightarrow \mathbb{R}$ is a convex and continuous function satisfying $f(1) = 0$. We also require $f(\cdot)$ is twice-differentiable. Let $\mathcal{F}[q]$ denote the energy functional $\mathbb{D}_f(\cdot\|p) : \mathcal{P}(\mathbb{R}^d) \rightarrow \mathbb{R}^+ \cup \{0\}$ for simplicity. We consider a curve $\mu_t : \mathbb{R}^+ \rightarrow \mathcal{P}(\mathbb{R}^d)$ and μ_t admits the density q_t . Let $\mathbf{v}_t = -\nabla \left(\frac{\delta \mathcal{F}}{\delta q_t}(q_t) \right) : \mathbb{R}^+ \rightarrow (\mathbb{R}^d \rightarrow \mathbb{R}^d)$ be the velocity vector field with $r_t(\mathbf{x}) = \frac{q_t(\mathbf{x})}{p(\mathbf{x})}$.

Definition 1 We call μ_t a variational gradient flow of the energy functional $\mathcal{F}[\cdot]$ governed by the velocity vector field \mathbf{v}_t if the following Vlasov-Fokker-Planck equation holds

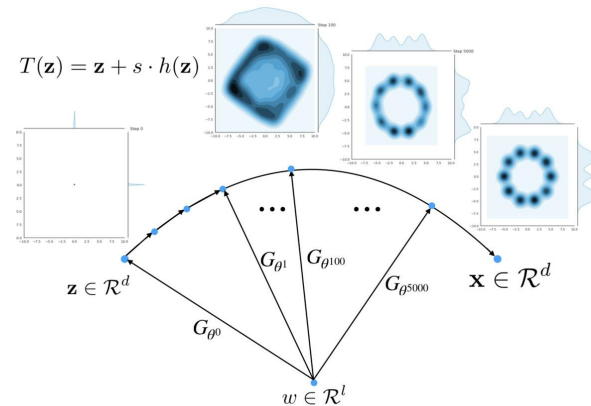
$$\frac{d}{dt} q_t = -\nabla \cdot (q_t \mathbf{v}_t) \quad \text{in } \mathbb{R}^+ \times \mathbb{R}^d. \quad (2)$$

Theorem 1 For any $\mathbf{g} \in \mathcal{H}(q_t)$, if the vanishing condition $\lim_{\|\mathbf{x}\| \rightarrow \infty} \|f'(r_t(\mathbf{x}))q_t(\mathbf{x})\mathbf{g}(\mathbf{x})\| = 0$ is satisfied, then

$$\left\langle \frac{\delta \mathcal{L}}{\delta \mathbf{h}}[0], \mathbf{g} \right\rangle_{\mathcal{H}(q_t)} = \langle f''(r_t) \nabla r_t, \mathbf{g} \rangle_{\mathcal{H}(q_t)}.$$

Theorem 2 The evolving distribution q_t under the infinitesimal pushforward map $\mathbb{T}_{s, \mathbf{v}_t}$ satisfies the Vlasov-Fokker-Planck equation (2).

Theorem 3 At the population level, the logD-trick GAN minimizes the “logD” divergence $\mathbb{D}_f(q(\mathbf{x})\|p(\mathbf{x}))$, with $f(u) = (u + 1) \log(u + 1) - 2 \log 2$, where $q(\mathbf{x})$ is the distribution of generated data.



Yuan, G., Yuling, J., Yang, W., Yao, W., Can, Y., & Shunkang, Z. (2019). Deep generative learning via variational gradient flow. *arXiv preprint arXiv:1901.08469*.


Reformulate GAN in function space



MNIST



Fashion-MNIST

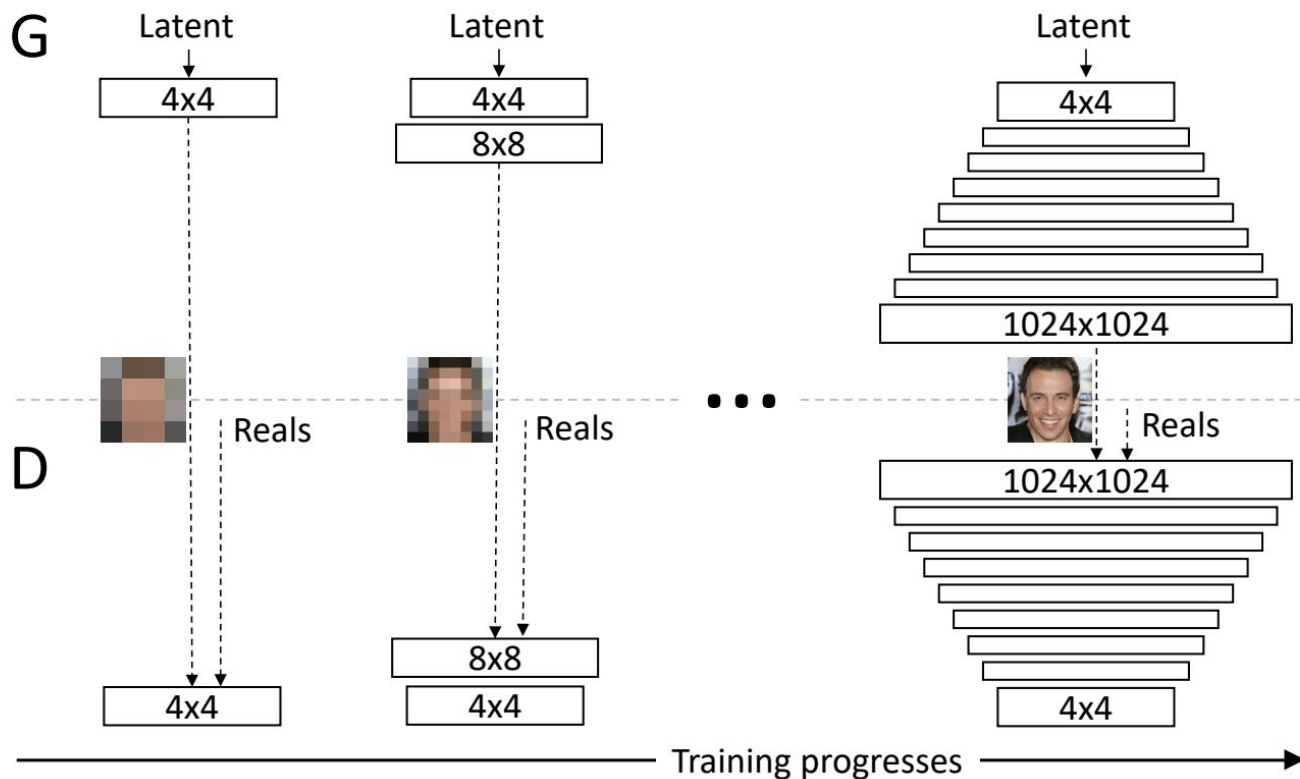


CIFAR10

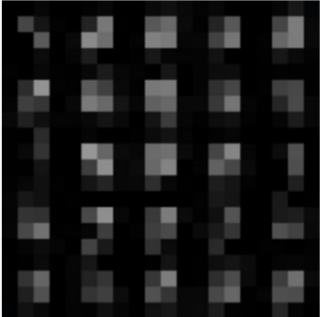
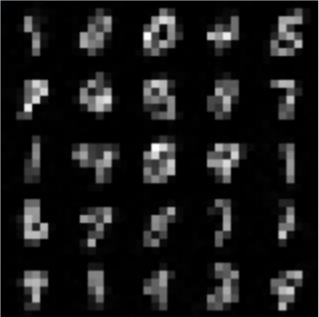


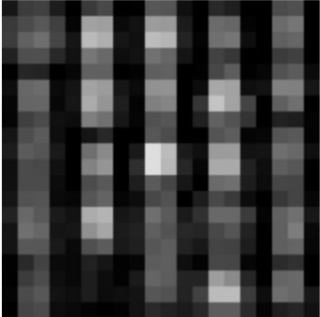
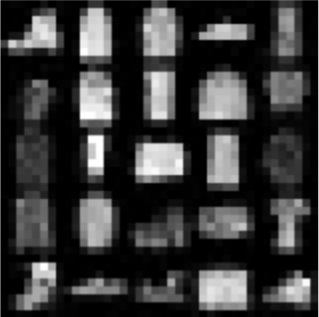




CelebA

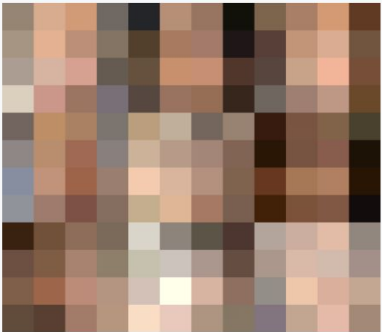





GAN with progressive training



GAN with progressive training

Resolution	4x4	8x8	16x16	32x32
MNIST				
Fashion-MNIST				

GAN with progressive training

Resolution	4x4	8x8	16x16
CelebA			
Resolution	32x32	64x64	128x128
CelebA			

Generative Adversarial Networks: A Review

CSIC 5011: Topological and Geometrical Data Reduction and Visualization

Spring 2020

Jose Vinicius de Miranda Cardoso, Yixin Men, Shunkang Zhang